

ME 759
High Performance Computing for Engineering Applications
Assignment 3
Due Thursday 02/12/2026 at 9:00 PM

Submit responses to all tasks which don't specify a file name to Canvas in a file called assignment3.{txt, docx, pdf, rtf, odt} (choose one of these formats). Submit all plots (if any) on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the **HW03** subdirectory on the **main** branch of your GitHub repo. Please use the name **HW03** exactly as shown here (both in terms of capitalization & name). The **HW03** subdirectory should have no subdirectories. For this assignment, your **HW03** folder should contain **task1.cu**, **task2.cu**, **task3.cu**, and **vscale.cu**.

All commands or code must work on *Euler* with only the **nvidia/cuda/13.0.0** module loaded. The commands may behave differently on your computer, so be sure to test on *Euler* before you submit.

Please submit clean code. Consider using a formatter like [clang-format](#).

IMPORTANT: Before you begin, copy any provided files from [Assignments/HW03](#) directory of the [ME759 Resource Repo](#). Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

Problem 1. Write a C++ program using CUDA in a file called **task1.cu** which computes the factorial of integers from 1 to 8, by launching a GPU kernel with 1 block and 8 threads. To that end:

- From the host, you need to allocate an array of 8 **ints** on the device called **dA**.
- You need to launch a kernel with 1 block having 8 threads.
- Each thread computes **(a+1)!**, where **a** is the thread's **threadIdx**, and writes the result in the **a**-th entry of the **dA** array.
- You need to copy back the data stored in the device array **dA** into a host array called **hA**.
- You need to print (from the host) the 8 values stored in the host array, one value per line
- Compile: **nvcc task1.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task1**
- Run (on Euler, use Slurm sbatch): **./task1**
- Expected output (showing only 4 out of the 8 lines expected):
1
2
6
24
- It is ok to publish your sbatch script on Piazza. For this assignment, you will need to ask Slurm to execute your program on a node that has at least one GPU card

Problem 2. Write a C++ program using CUDA in a file called **task2.cu** which does the following:

- From the host, allocates an array of 16 **ints** on the device called **dA**.
- Launches a kernel with 2 blocks, each block having 8 threads.
- Each thread computes **ax+y** and writes the result in one distinct entry of the **dA** array (take a look below at the expected output of the program to figure out which entry of the array a thread needs to write to). Here,
 - **x** is the thread's **threadIdx**;
 - **y** is the thread's **blockIdx**;

- `a` is an integer argument that the kernel takes (so all threads use the same `a`). You need to generate `a` randomly and then call the kernel with it. It is up to you how you generate this random number, one possible approach is described here [BestPractice](#).
- Copies back the data stored in the device array `dA` into a host array called `hA`.
- Prints (from the host) the 16 sequential values stored in the host array separated by a single space each.

How to go about it, and what the expected output looks like:

- Compile: `nvcc task2.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task2`
- Run (on Euler, use Slurm sbatch): `./task2`
- Expected output (followed by newline; yours could be different depending on the random number generation): `0 10 20 30 40 50 60 70 1 11 21 31 41 51 61 71`

- Problem 3.** a) Implement in a file called `vscale.cu`, the `vscale` kernel function as declared and described in `vscale.cuh`. This function should take in two arrays, `a` and `b`, and do an element-wise multiplication of the two arrays: $b_i = a_i \cdot b_i$. In the process, `b` gets overwritten. Each thread should do at most one of the multiplication operations.

Example:

$$\mathbf{a} = [-5.0, 2.0, 1.5], \quad \mathbf{b} = [0.8, 0.3, 0.6], \quad \mathbf{n} = 3$$

The resulting `b` array is:

$$\mathbf{b} = [-4.0, 0.6, 0.9].$$

- b) Write a file `task3.cu` which does the following:

- Creates two arrays of length `n` filled by random numbers¹ where `n` is read from the first command line argument. The range of values for array `a` is $[-10.0, 10.0]$, whereas the range of values for array `b` is $[0.0, 1.0]$.
- Calls your `vscale` kernel with a 1D execution configuration that uses 512 threads per block.
- Prints the amount of time taken to execute the kernel in *milliseconds* using CUDA events².
- Prints the first element of the resulting array.
- Prints the last element of the resulting array.

How to go about it, and what the expected output looks like:

- Compile: `nvcc task3.cu vscale.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task3`
- Run (on Euler, use Slurm sbatch): `./task3 n`
- Example expected output (followed by newline):


```
0.012
      1.3
      2.3
```
- We will only test with positive integers for `n`.
- c) On an Euler *compute node*, run `task3` for each value $n = 2^{10}, 2^{11}, \dots, 2^{29}$ and generate a plot `task3.pdf` which plots the time taken by your `vscale` as a function of `n`. Overlay another plot which shows the scaling results when using 16 threads per block.

¹Details about random number generation can be found in [random_numbers.md](#).

²Recall the GPU timing section of the document [timing.md](#).